

6.035

Memory Optimization

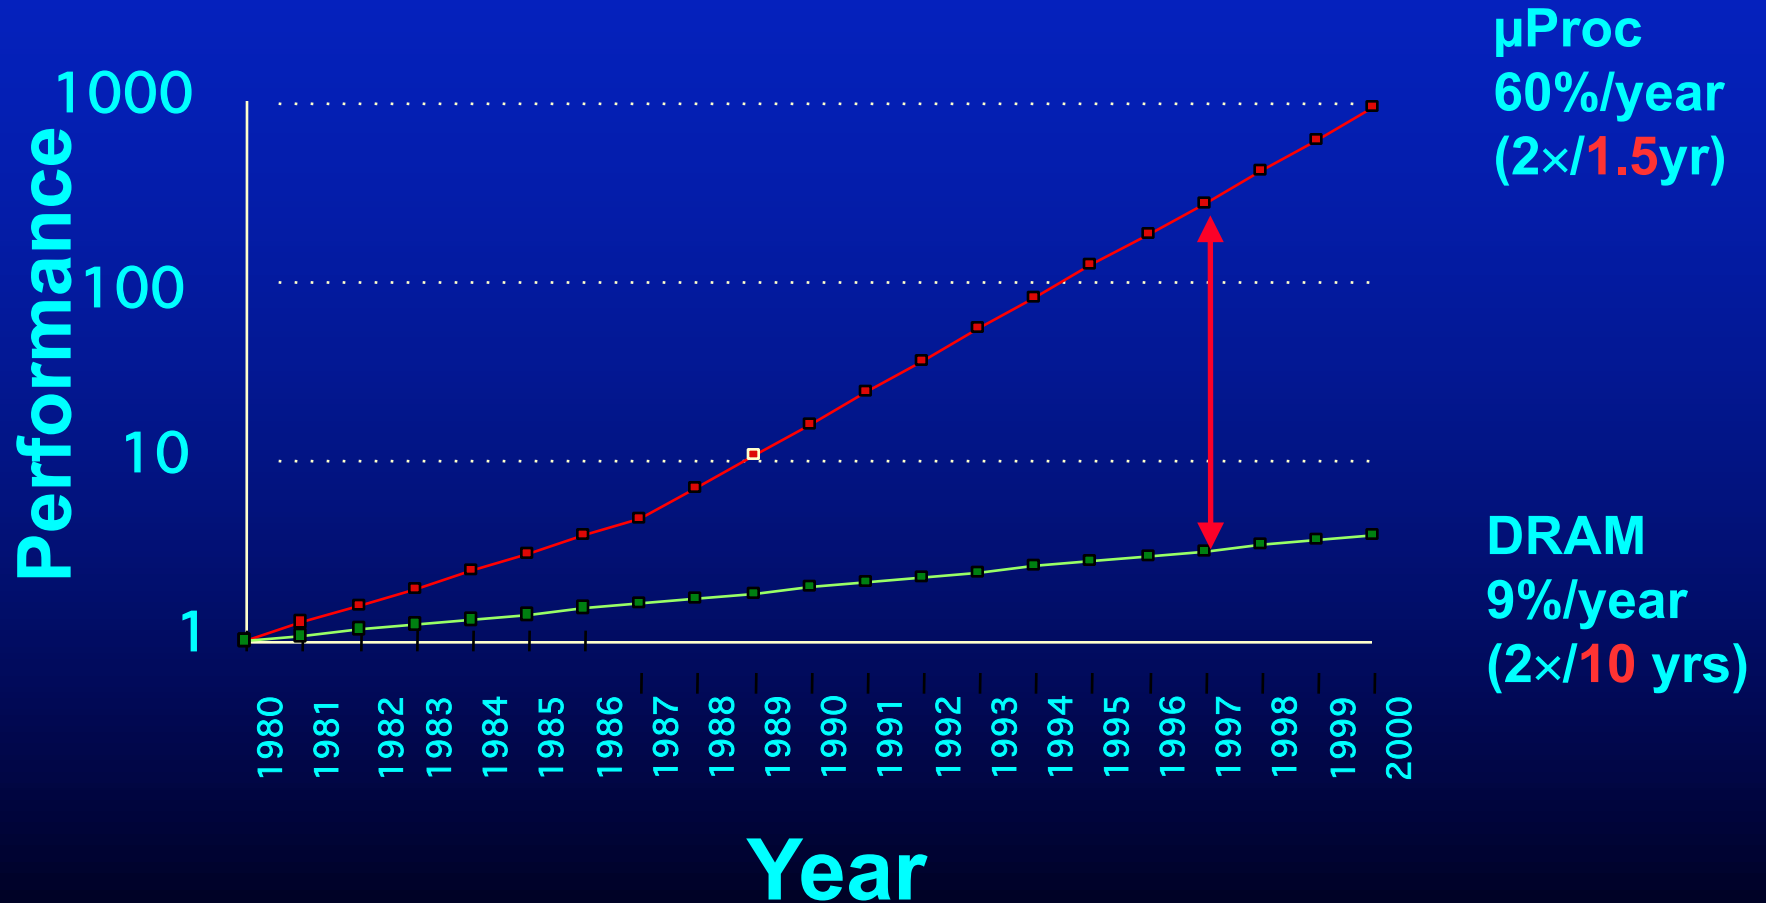
Outline

- Issues with the Memory System
- Loop Transformations
- Data Transformations
- Prefetching
- Alias Analysis

Memory Hierarchy

1 - 2 ns	Registers	32 – 512 B
3 - 10 ns	L1 Private Cache	16 – 128 KB
8 - 30 ns	L2/L3 Shared Cache	1 – 16 MB
60 - 250 ns	Main Memory (DRAM)	1 GB – 128 GB
5 - 20 ms	Permanent Storage (Hard Disk)	250 GB – 4 TB

Processor-Memory Gap



Cache Architecture

		Pentium D	Core Duo	Core 2 Duo	Athlon 64
L1 code (per core)	size	12 K uops	32 KB	32 KB	64 KB
	associativity	8 way	8 way	8 way	2 way
	Line size	64 bytes	64 bytes	64 bytes	64 bytes
L1 data (per core)	size	16 KB	32 KB	32 KB	64 KB
	associativity	8 way	8 way	8 way	8 way
	Line size	64 bytes	64 bytes	64 bytes	64 bytes
L1 to L2	Latency	4 cycles	3 cycles	3 cycles	3 cycles
L2 shared	size	4 MB	4 MB	4 MB	1 MB
	associativity	8 way	8 way	16 way	16 way
	Line size	64 bytes	64 bytes	64 bytes	64 bytes
L2 to L3(off)	Latency	31 cycles	14 cycles	14 cycles	20 cycles

Cache Misses

- Cold misses
 - First time a data is accessed
- Capacity misses
 - Data got evicted between accesses because a lot of other data (more than the cache size) was accessed
- Conflict misses
 - Data got evicted because a subsequent access fell on the same cache line (due to associativity)
- True sharing misses (multicores)
 - Another processor accessed the data between the accesses
- False sharing misses (multicores)
 - Another processor accessed different data in the same cache line between the accesses

Data Reuse

- Temporal Reuse

- A given reference accesses the same location in multiple iterations

```
for i = 0 to N
  for j = 0 to N
    A[j] =
```

- Spatial Reuse

- Accesses to different locations within the same cache line

```
for i = 0 to N
  for j = 0 to N
    B[i, j] =
```

- Group Reuse

- Multiple references access the same location

```
for i = 0 to N
  A[i] = A[i-1] + 1
```

Outline

- Issues with the Memory System
- **Loop Transformations**
- Data Transformations
- Prefetching
- Alias Analysis

Matrix Multiply

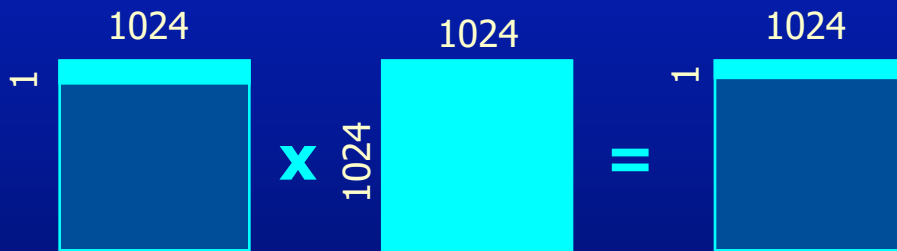
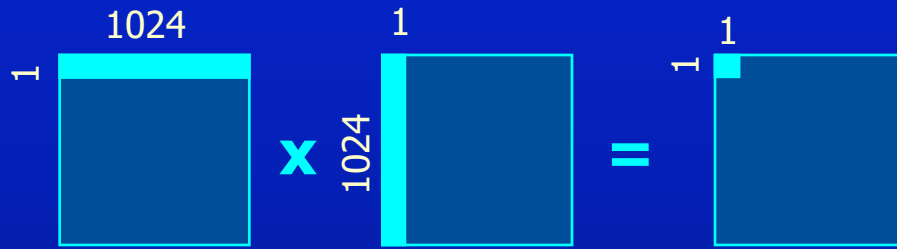
for i = 1 to n

 for j = 1 to n

 for k = 1 to n

$c[i,j] += a[i,k] * b[k,j]$

Example: Matrix Multiply



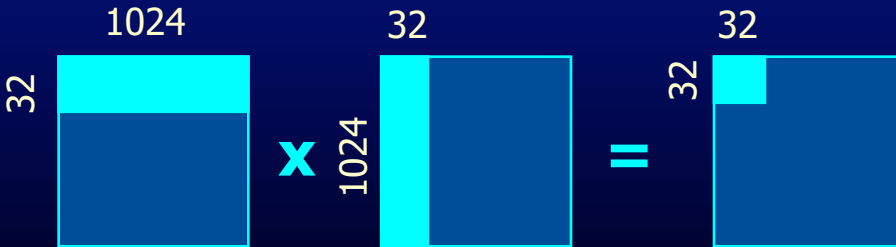
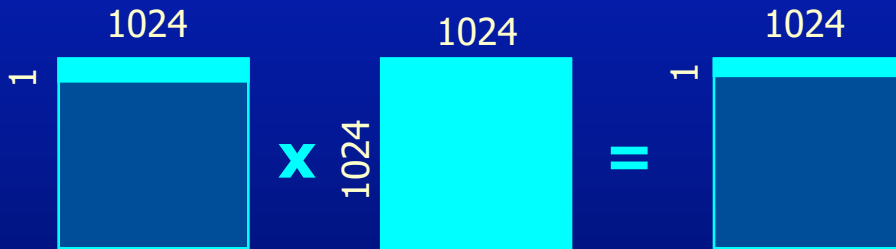
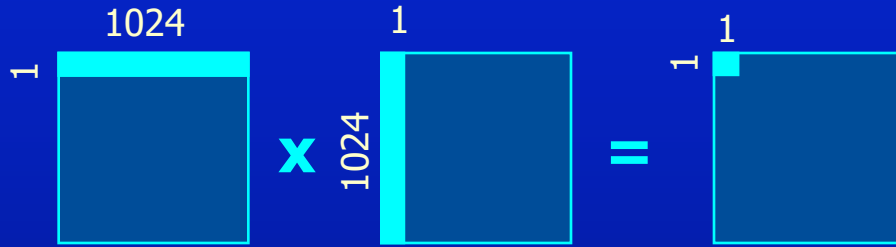
Data Accessed

1,050,624

Matrix Multiply

```
for i0 = 1 to n step b
  for j0 = 1 to n step b
    for k0 = 1 to n step b
      for i = i0 to min(i0+b-1, n)
        for j = j0 to min(j0+b-1, n)
          for k = k0 to min(k0+b-1, n)
            c[i,j] += a[i,k]*b[k,j]
```

Example: Matrix Multiply



Data Accessed

1,050,624

66,560

Loop Transformations

- Transform the iteration space to reduce the number of misses
- Reuse distance – For a given access, number of other data items accessed before that data is accessed again
- Reuse distance $>$ cache size
 - Data is spilled between accesses

Divide and Conquer Matrix Multiply

$$\begin{array}{|c|c|} \hline A & B \\ \hline C & D \\ \hline \end{array} \times \begin{array}{|c|c|} \hline E & F \\ \hline G & H \\ \hline \end{array} = \begin{array}{|c|c|} \hline AE+BG & AF+BH \\ \hline CE+DG & CF+DH \\ \hline \end{array}$$

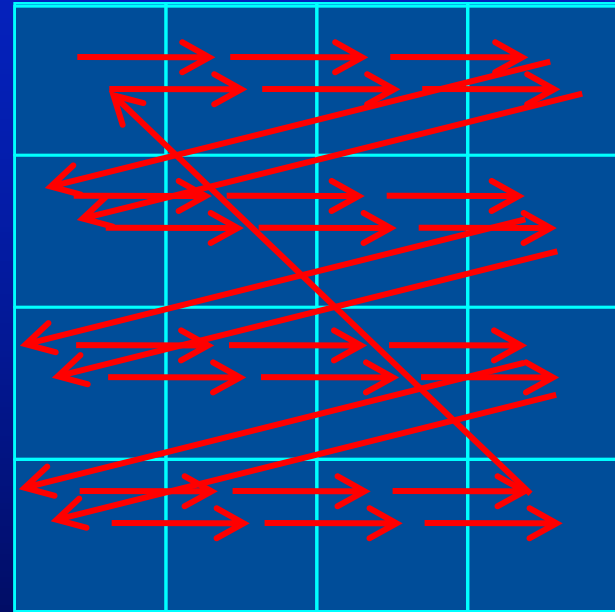
Loop Transformations

```
for i = 0 to N
  for j = 0 to N
    for k = 0 to N
      A[k,j]
```

Reuse distance = N^2

If Cache size < 16 doubles?

A lot of capacity misses

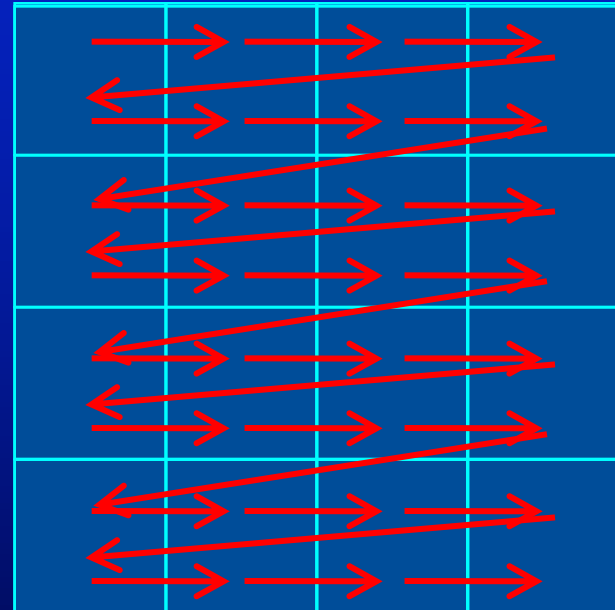


Loop Transformations

```
for i = 0 to N
  for j = 0 to N
    for k = 0 to N
      A[k,j]
```

Loop Interchange

```
for j = 0 to N
  for i = 0 to N
    for k = 0 to N
      A[k,j]
```



Loop Transformations

```
for j = 0 to N
  for i = 0 to N
    for k = 0 to N
      A[k,j]
```

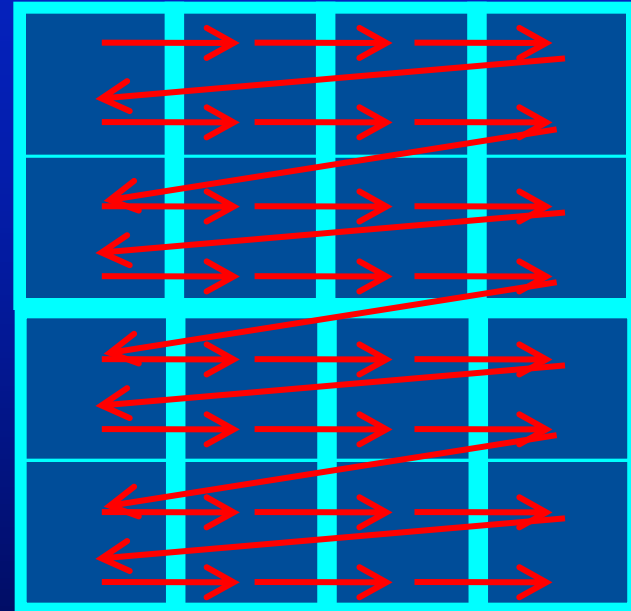
Cache line size > data size

Cache line size = L

Reuse distance = LN

If cache size < 8 doubles?

Again a lot of capacity misses

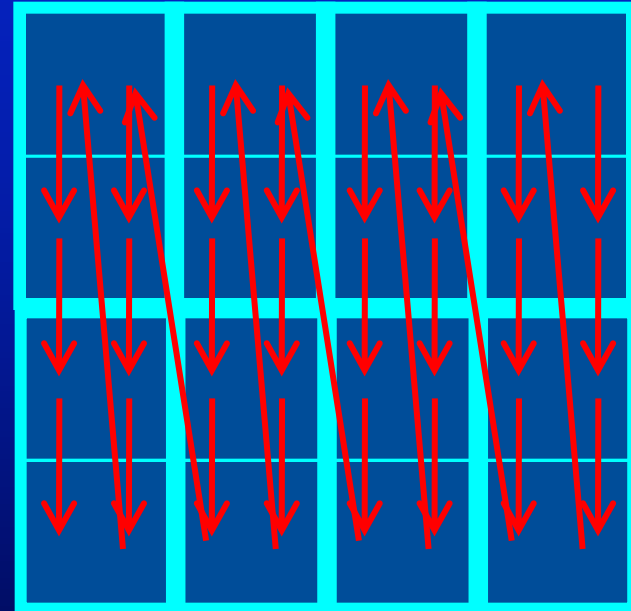


Loop Transformations

```
for j = 0 to N
  for i = 0 to N
    for k = 0 to N
      A[k,j]
```

Loop Interchange

```
for k = 0 to N
  for i = 0 to N
    for j = 0 to N
      A[k,j]
```



Loop Transformations

for i = 0 to N

 for j = 0 to N

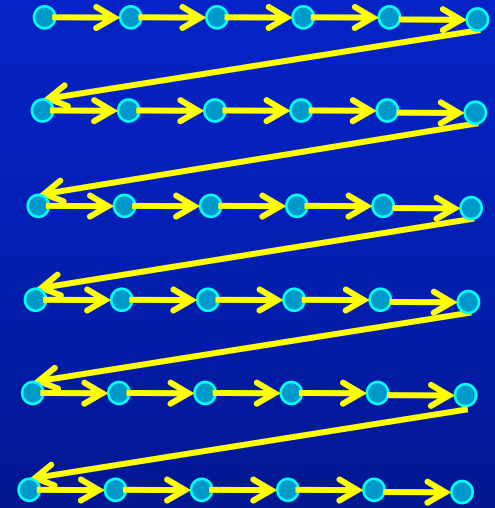
 for k = 0 to N

$A[i,j] = A[i,j] + B[i,k] + C[k,j]$

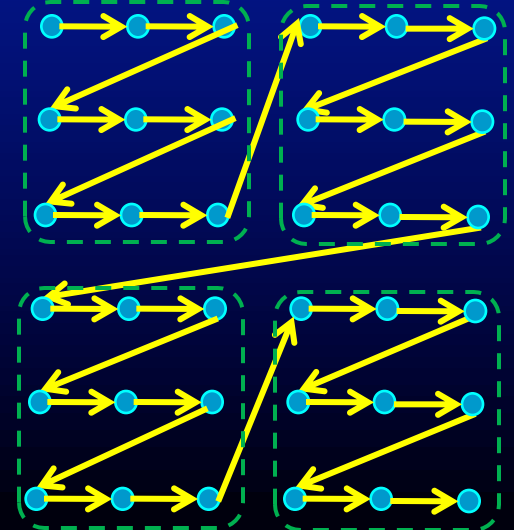
- No matter what loop transformation you do one array access has to traverse the full array multiple times

Loop Tiling

```
for i = 0 to N  
  for j = 0 to N
```



```
for ii = 0 to ceil(N/b)  
  for jj = 0 to ceil(N/b)  
    for i = b*ii to min(b*ii+b-1, N)  
      for j = b*jj to min(b*jj+b-1, N)
```

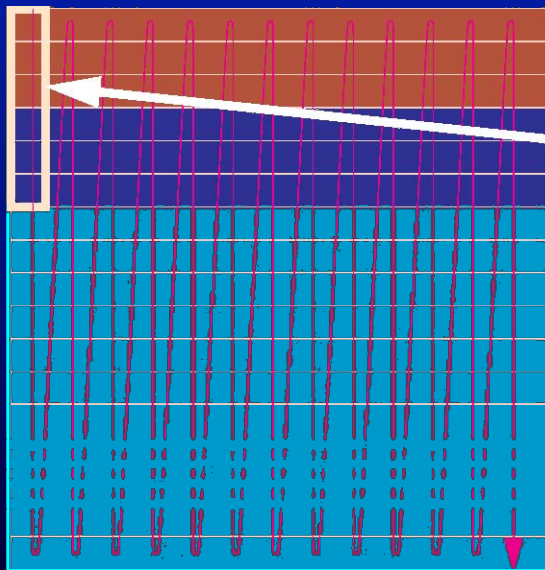


Outline

- Issues with the Memory System
- Loop Transformations
- **Data Transformations**
- Prefetching
- Alias Analysis

False Sharing Misses

```
for J =  
  forall I =  
    X(I, J) = ...
```

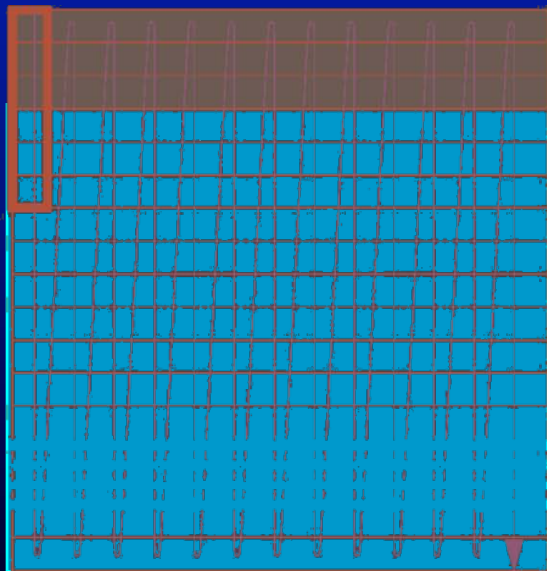


Cache Lines

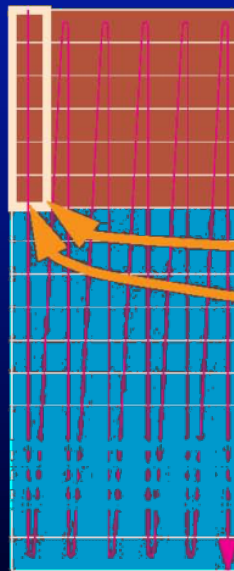
Array X

Conflict Misses

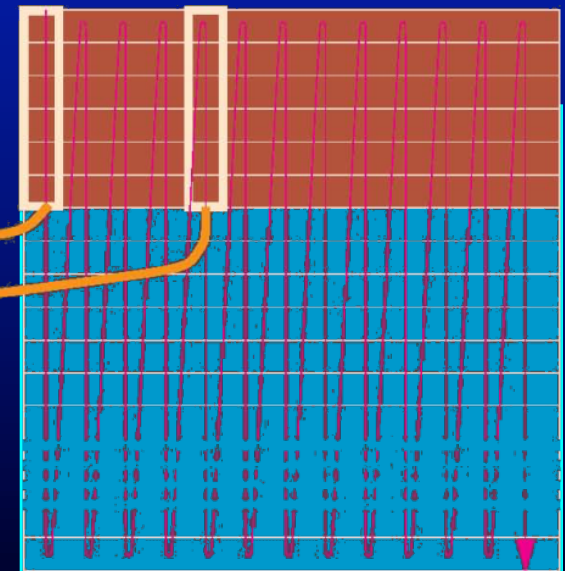
```
for J =  
  forall I =  
    X(I, J) = ...
```



Array X



Cache



Memory

Data Transformations

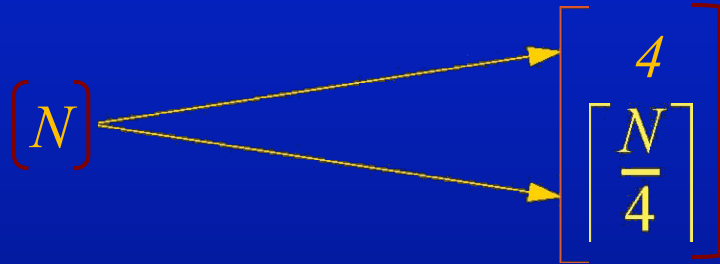
- Similar to loop transformations
- All the accesses have to be updated
 - Whole program analysis is required

Strip-Mining

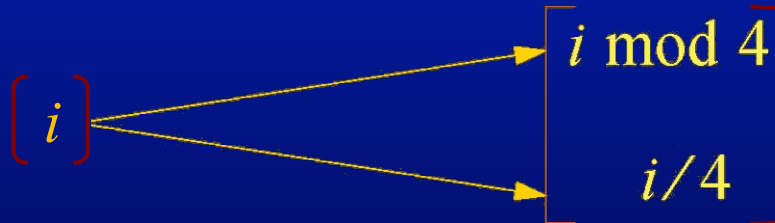
Create two dims from one

With blocksize=4

Storage
Declaration



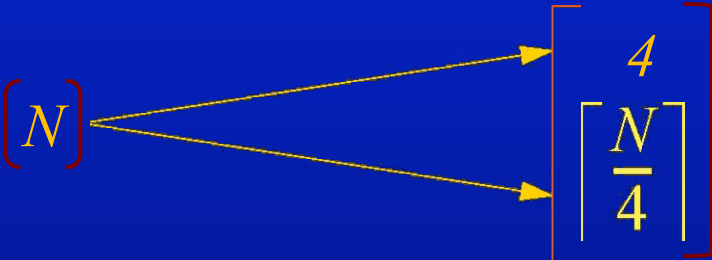
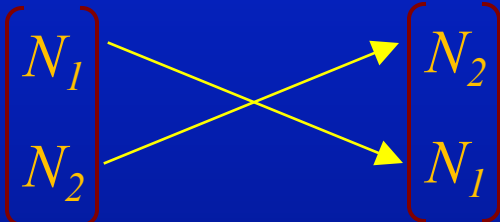
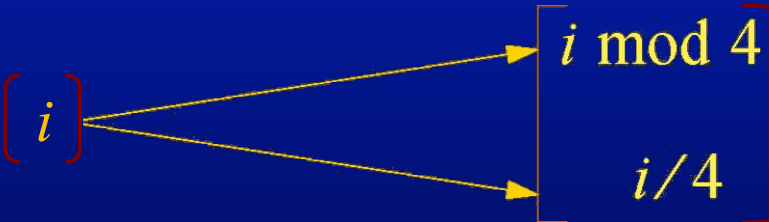
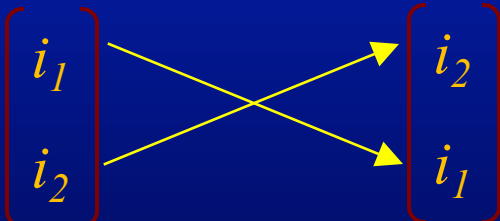
Array
Access



Memory
Layout

0	1	2	3	4	5	6	7
0	1	2	3	4	5	6	7

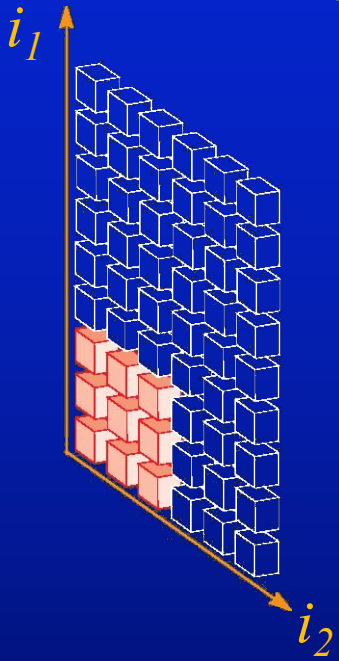
0	1	2	3	4	5	6	7
$0,0$	$1,0$	$2,0$	$3,0$	$0,1$	$1,1$	$2,1$	$3,1$

	<div>Strip-Minding</div> <div>Create two dims from one</div> <div>With blocksize=4</div>	<div>Permutation</div> <div>Change memory layout</div> <div>With permutation matrix $\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$</div>																																																																																
Storage Declaration																																																																																		
Array Access																																																																																		
Memory Layout	<table><tr><td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td></tr><tr><td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td></tr></table> <table><tr><td>0</td><td>1</td><td>2</td><td>3</td><td>4</td><td>5</td><td>6</td><td>7</td></tr><tr><td>0,0</td><td>1,0</td><td>2,0</td><td>3,0</td><td>0,1</td><td>1,1</td><td>2,1</td><td>3,1</td></tr></table>	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	0,0	1,0	2,0	3,0	0,1	1,1	2,1	3,1	<table><tr><td>0</td><td>4</td><td>8</td></tr><tr><td>0,0</td><td>0,1</td><td>0,2</td></tr><tr><td>1</td><td>5</td><td>9</td></tr><tr><td>1,0</td><td>1,1</td><td>1,2</td></tr><tr><td>2</td><td>6</td><td>10</td></tr><tr><td>2,0</td><td>2,1</td><td>2,2</td></tr><tr><td>3</td><td>7</td><td>11</td></tr><tr><td>3,0</td><td>3,1</td><td>3,2</td></tr></table> <table><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>0,0</td><td>1,0</td><td>2,0</td></tr><tr><td>3</td><td>4</td><td>5</td></tr><tr><td>0,1</td><td>1,1</td><td>2,1</td></tr><tr><td>6</td><td>7</td><td>8</td></tr><tr><td>0,2</td><td>1,2</td><td>2,2</td></tr><tr><td>9</td><td>10</td><td>11</td></tr><tr><td>0,3</td><td>1,3</td><td>2,3</td></tr></table>	0	4	8	0,0	0,1	0,2	1	5	9	1,0	1,1	1,2	2	6	10	2,0	2,1	2,2	3	7	11	3,0	3,1	3,2	0	1	2	0,0	1,0	2,0	3	4	5	0,1	1,1	2,1	6	7	8	0,2	1,2	2,2	9	10	11	0,3	1,3	2,3
0	1	2	3	4	5	6	7																																																																											
0	1	2	3	4	5	6	7																																																																											
0	1	2	3	4	5	6	7																																																																											
0,0	1,0	2,0	3,0	0,1	1,1	2,1	3,1																																																																											
0	4	8																																																																																
0,0	0,1	0,2																																																																																
1	5	9																																																																																
1,0	1,1	1,2																																																																																
2	6	10																																																																																
2,0	2,1	2,2																																																																																
3	7	11																																																																																
3,0	3,1	3,2																																																																																
0	1	2																																																																																
0,0	1,0	2,0																																																																																
3	4	5																																																																																
0,1	1,1	2,1																																																																																
6	7	8																																																																																
0,2	1,2	2,2																																																																																
9	10	11																																																																																
0,3	1,3	2,3																																																																																

Data Transformation Algorithm

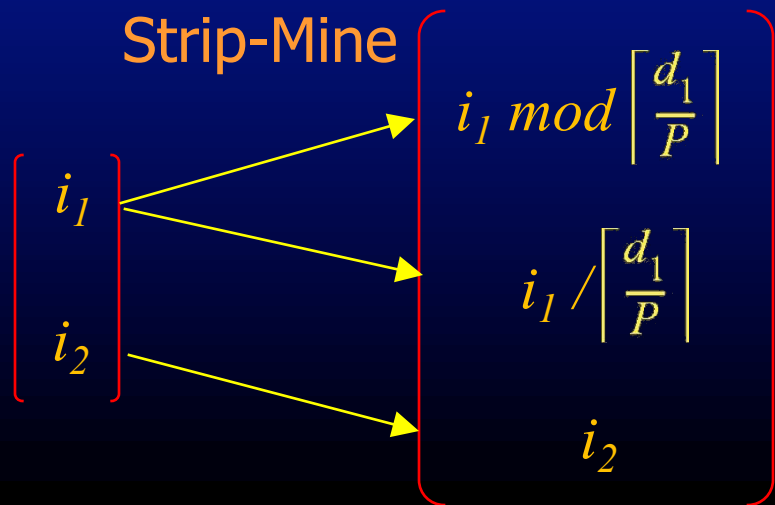
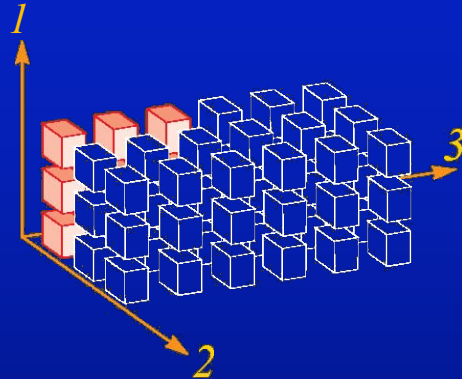
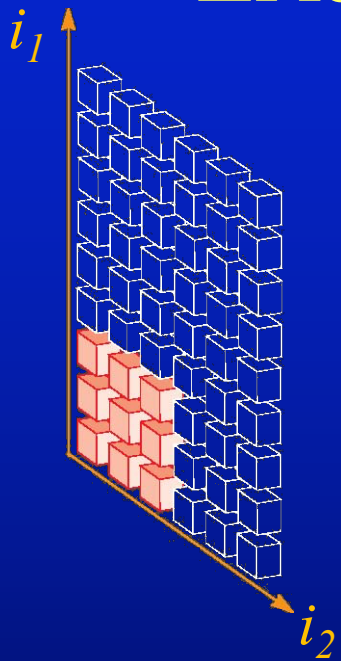
- Rearrange data: Each processor's data is contiguous
- Use data decomposition
 - *, block, cyclic, block-cyclic
- Transform each dimension according to the decomposition
- Use a combination of strip-mining and permutation primitives

Example I: (Block, Block)

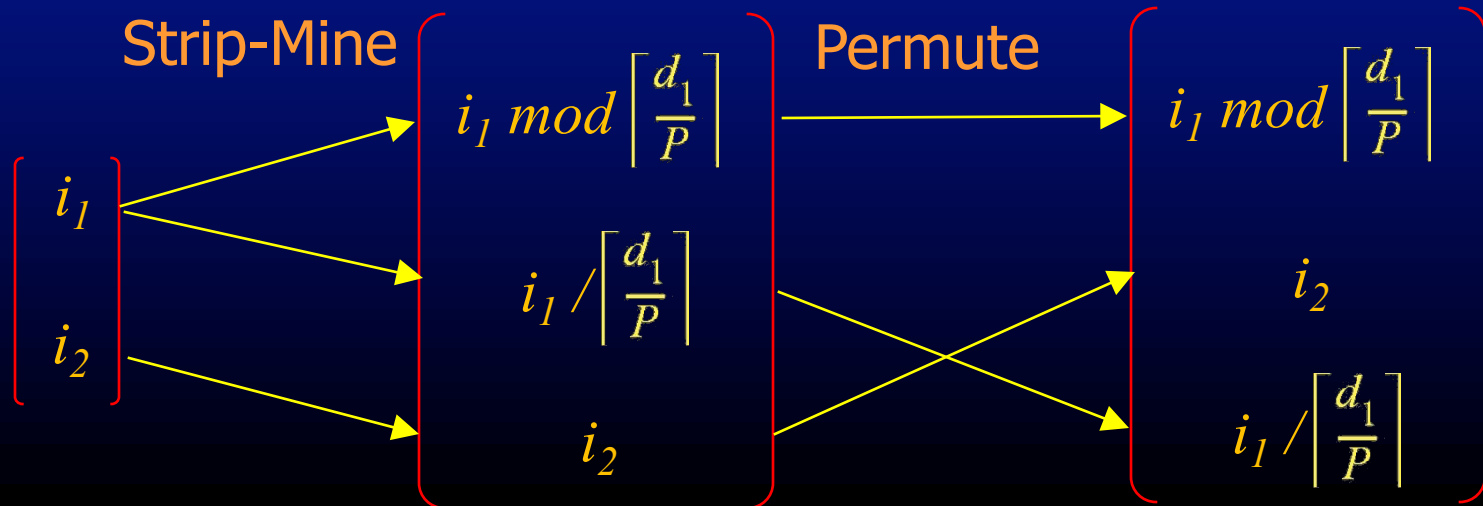
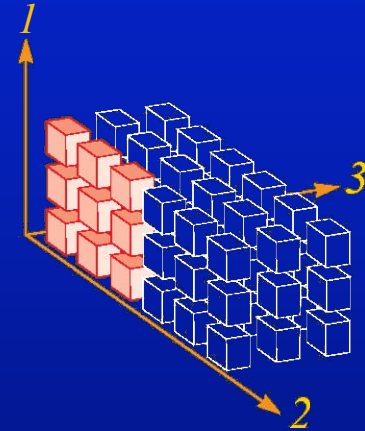
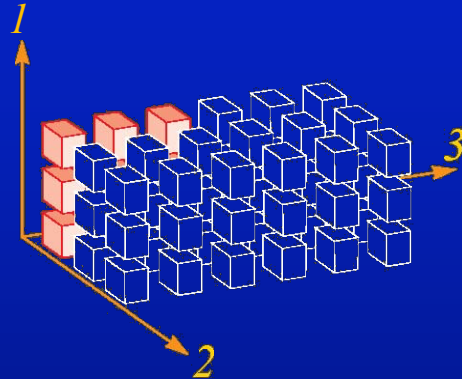
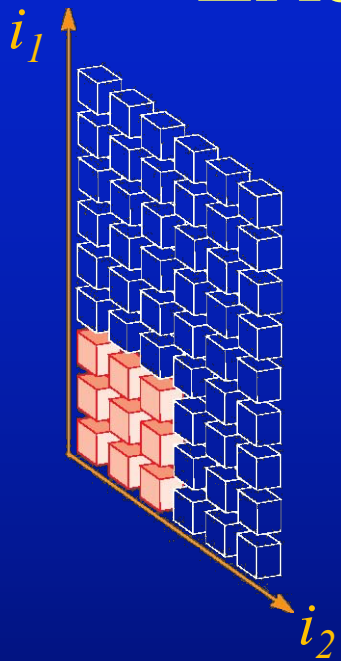


$$\begin{bmatrix} i_1 \\ i_2 \end{bmatrix}$$

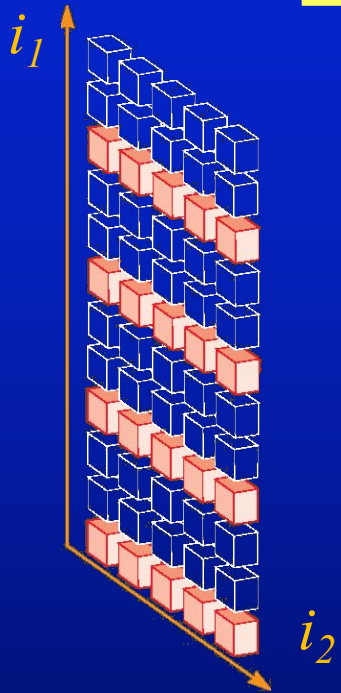
Example I: (Block, Block)



Example I: (Block, Block)

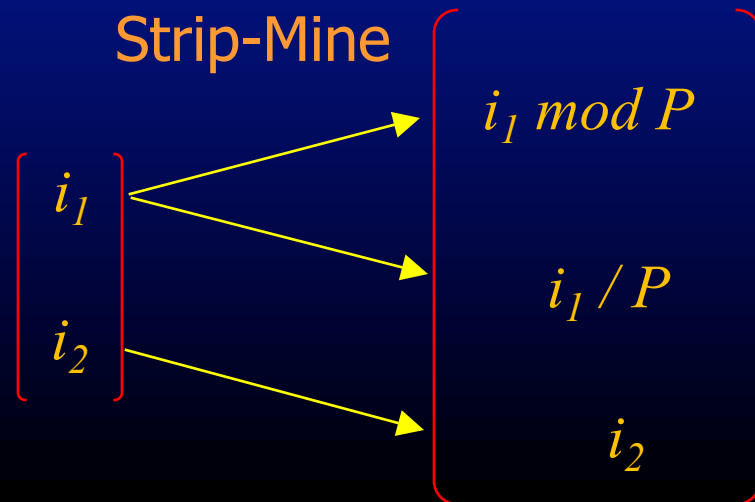
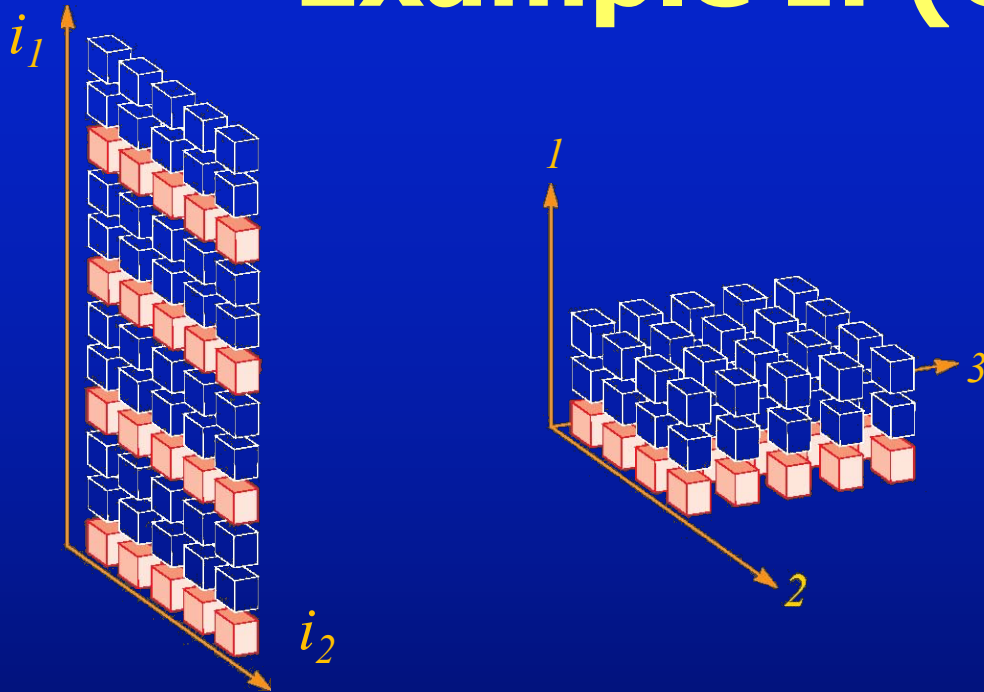


Example I: (Cyclic, *)

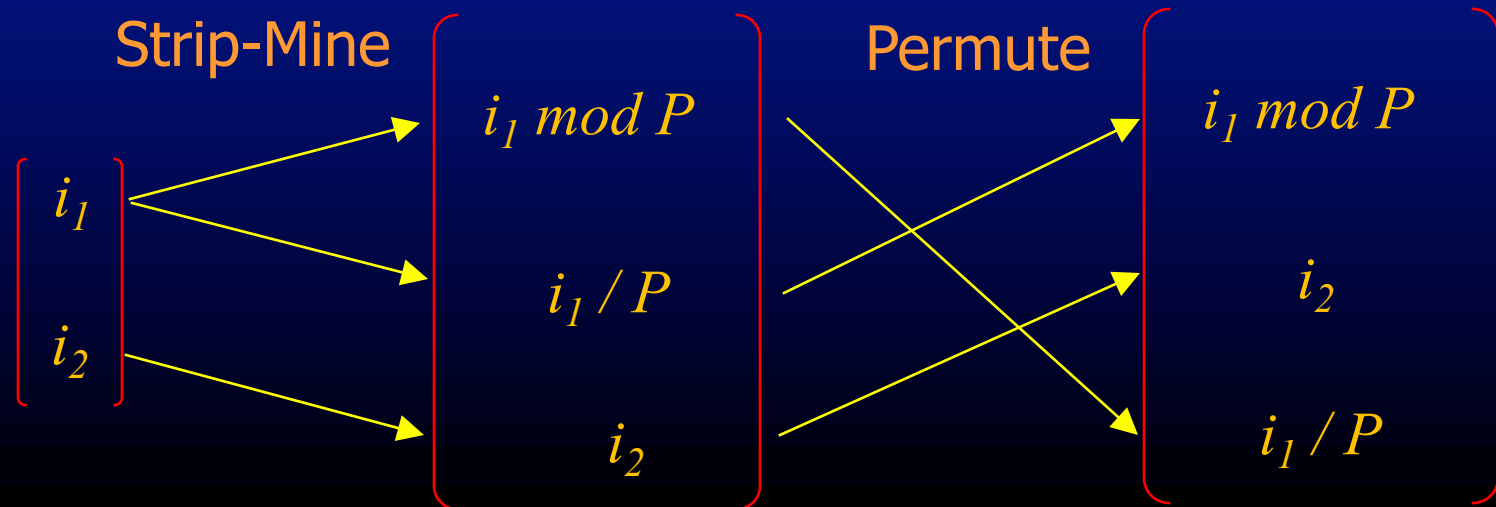
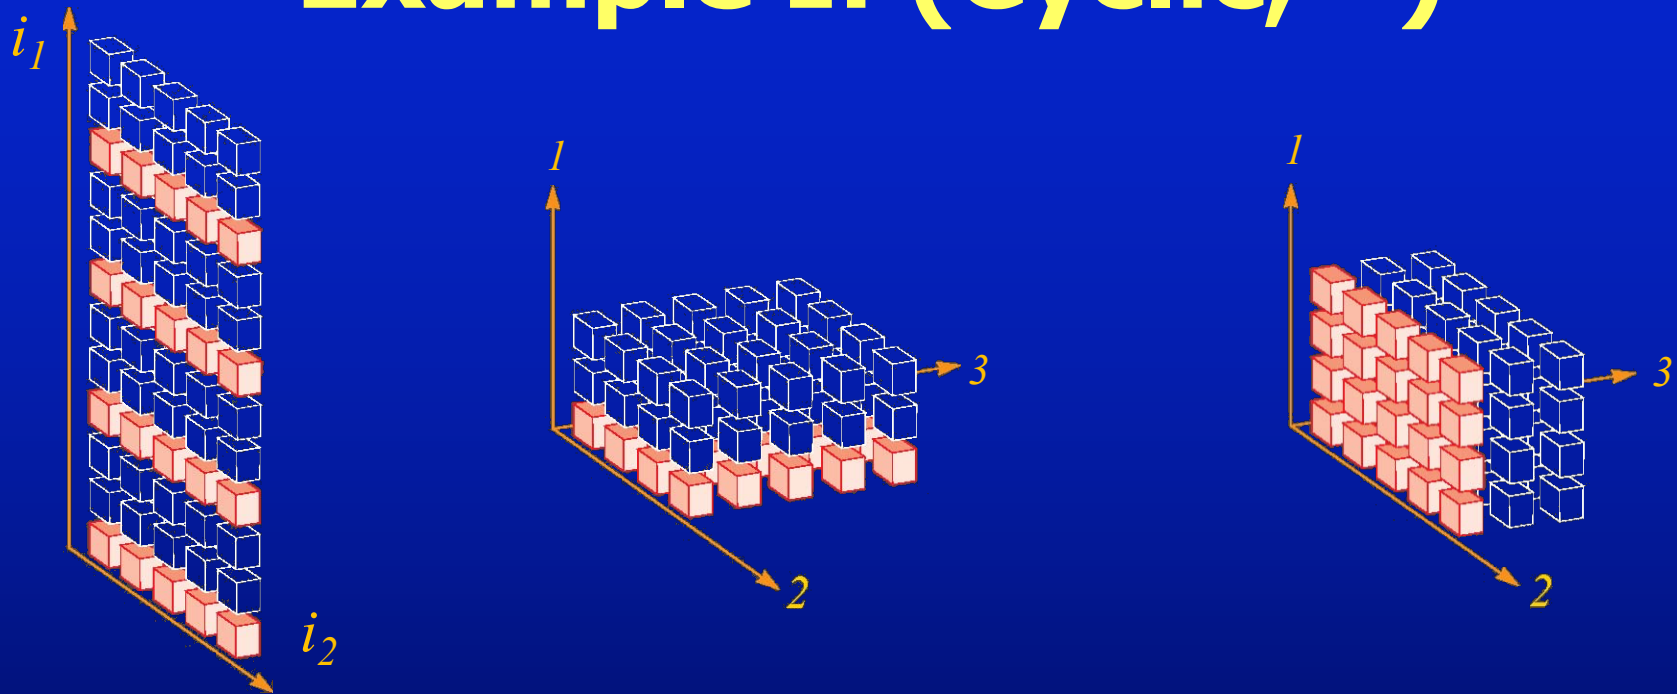


$$\begin{bmatrix} i_1 \\ i_2 \end{bmatrix}$$

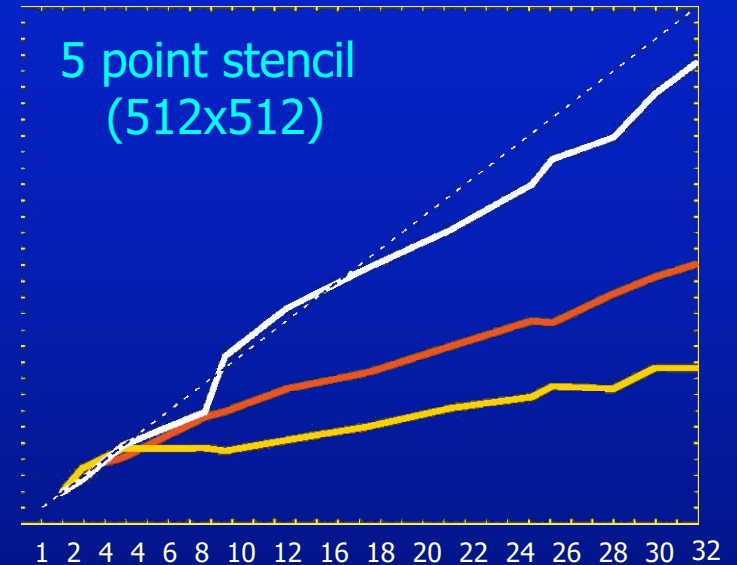
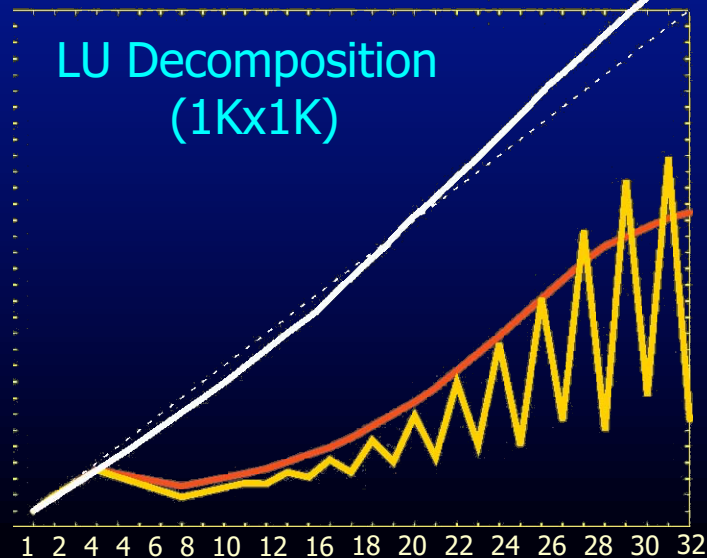
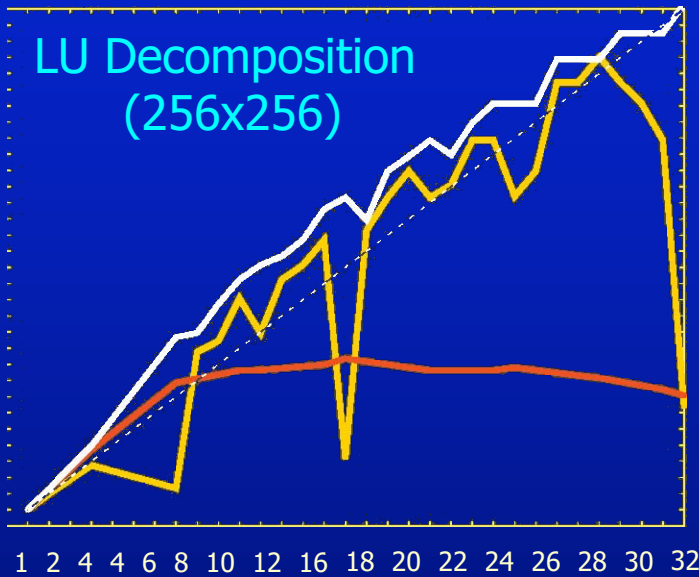
Example I: (Cyclic, *)



Example I: (Cyclic, *)



Performance



- Parallelizing outer loop
- Best computation placement
- + data transformations

Optimizations

- Modulo and division operations in the index calculation
 - Very high overhead
- Use standard techniques
 - Loop invariant removal, CSE
 - Strength reduction exploiting properties of modulo and division
 - Use knowledge about the program

Outline

- Issues with the Memory System
- Loop Transformations
- Data Transformations
- **Prefetching**
- Alias Analysis

Prefetching

- Cache miss stalls the processor for hundreds of cycles
 - Start fetching the data early so it'll be available when needed
- Pros
 - Reduction of cache misses → increased performance
- Cons
 - Prefetch contents for fetch bandwidth
 - Solution: Hardware only issue prefetches on unused bandwidth
 - Evicts a data item that may be used
 - Solution: Don't prefetch too early
 - Prefetch is still pending when the memory is accessed
 - Solution: Don't prefetch too late
 - Prefetch data is never used
 - Solution: Prefetch only data guaranteed to be used
 - Too many prefetch instructions
 - Prefetch only if access is going to miss in the cache

Prefetching

- Compiler inserted
 - Use reuse analysis to identify misses
 - Partition the program and insert prefetches
- Run ahead thread (helper threads)
 - Create a separate thread that runs ahead of the main thread
 - Runahead only does computation needed for control-flow and address calculations
 - Runahead performs data (pre)fetches

Outline

- Issues with the Memory System
- Loop Transformations
- Data Transformations
- Prefetching
- **Alias Analysis**

Alias Analysis

- Aliases destroy local reasoning
 - Simple, local transformations require global reasoning in the presence of aliases
 - A critical issue in pointer-heavy code
 - This problem is even worse for multithreaded programs
- Two solutions
 - Alias analysis
 - Tools to tell us the potential aliases
 - Change the programming language
 - Languages have no facilities for talking about aliases
 - Want to make local reasoning possible

Aliases

- Definition

*Two pointers that point to the same location are **aliases***

- Example

$Y = \&Z$

$X = Y$

$*X = 3$ */* changes the value of $*Y$ */*

Example

```
foo(int * A, int * B, int * C, int N)
  for i = 0 to N-1
    A[i] = A[i] + B[i] + C[i]
```

- Is this loop parallel?
- Depends

```
int X[1000];
int Y[1000];
int Z[1000]
foo(X, Y, Z, 1000);
```

```
int X[1000];
foo(&X[1], &X[0], &X[2], 998);
```

Points-To Analysis

- Consider:

$P = \&Q$

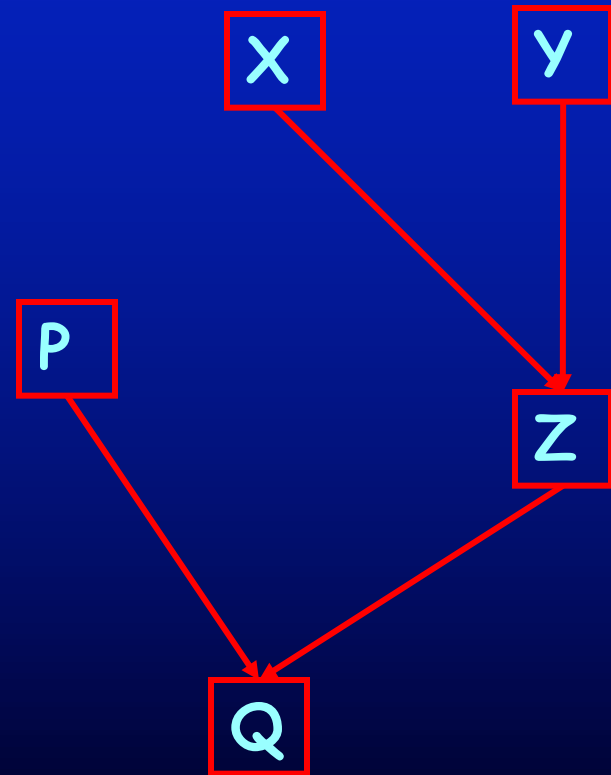
$Y = \&Z$

$X = Y$

$*X = P$

- Informally:

- P can point to Q
- Y can point to Z
- X can point to Z
- Z can point to Q



Points-To Relations

- A graph
 - Nodes are program *names*
 - Edge (x,y) says x may point to y
- Finite set of names
 - Implies each name represents many heap cells
 - Correctness: If $*x = y$ in any state of any execution, then (x,y) is an edge in the points-to graph

Sensitivity

- *Context sensitivity*
 - Separate different uses of functions
 - Different is the key – if the analysis think the input is the same, reuse the old results
- *Flow sensitivity*
 - For insensitivity makes any permutation of program statements gives same result
 - Flow sensitive is similar to data-flow analysis

Conclusion

- Memory systems are designed to give a huge performance boost for “normal” operations
- The performance gap between good and bad memory usage is huge
- Programs analyses and transformations are needed
- Can off-load this task to the compiler