

(9)

LECTURE 4  
INDEXING AND ACCESS METHODS  
PHYSICAL QUERY PLANS

Summary of Execution Models:

- Iterator Model (Volcano, pipeline model)
  - ↳ void open()
  - ↳ Tuple next()
  - ↳ void close()
  - ↳ "Pulls" tuples from the top
  - ↳ Easy to control intermediates. Easy stop (i.e. LIMIT)
- Batch-at-a-time
  - ↳ Reduce function calls (important for in-mem DBs)
- Bottom-up
  - ↳ Better locality. More memory required.

Logical and Physical (Execution) Query Plans:

- Logical QP uses relational algebra operators and extensions, and assume a way of reading data.
  - ↳ Describe the order in which operators are applied to execute a query.
  - ↳ They don't describe how each op. is implemented, or how data is precisely accessed.
- Both op. implementation and how data is accessed (access methods) are crucial for achieving good execution time/performance.
  - ↳ The role of the query optimizer is to find a good physical plan, op. impl. <sup>given</sup> and access methods.

Outline of this lecture:

- Overview of op. impl. and storage
- Example. Building query cost intuition
- Access Methods

Overview Op. Implementation and Storage:

- Nested loop join,  $R \bowtie S$ 

```
for r in R:
  for s in S:
    if predicate(r,s):
      output join(r,s)
```

• Hash Join

- ↳ Build hash table (in-mem) for the smallest relation.
- ↳ Probe tuples from the second (by scanning it).

→ Many more implementations. One lecture on 'Join Algos'.

Storage: (assume magnetic disk or SSD)

- Records (or tuples) are stored in "pages"
- Pages are part of HeapFiles.
- All records are unordered. Pages must keep track of which "slots" are free and occupied
- Pages are sized so it's efficient to read them from disk and write them to it.
- Pages are cached in the BufferPool, which works as a 'cache' using the main memory of the server.
- Why not storing tuples ordered in disk? What happens when we write a new one?
- Access Methods are strategies to read tuples from disk, knowing how storage is organized.

Performance Engineering. Building Intuition on Query Plan cost:

- CPU cost (# instructions / unit time)
- IO cost (# pages read / unit time)
- Random IO (page read + seek)

$$(1 \text{ GHz} \equiv 1 \text{ B instr/s}), \quad 1 \text{ ns/instr.}$$

$$100 \text{ MB/s} \equiv 10 \text{ ns/byte}$$

$$10 \text{ ms/seek} \equiv 100 \text{ seeks/sec}$$

triggers seek.

\* It's possible to execute around 10M instr/seek. Beware random IO

- What cost dominates in database execution?

- ↳ Depends on query workload (queries and data).
- ↳ Depends on hardware characteristics.
- ↳ Depends on database design (on-disk, in-memory).
- ↳ Depends on system implementation.

- When designing a system / trying to understand one, we must be able to make sense of the performance we observe.

Example:

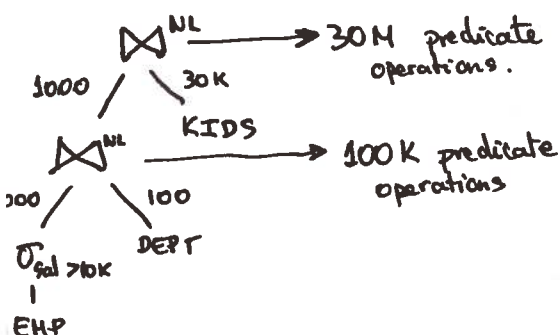
- 100 tuples/page
- 10 KB/page
- 10 pages RAM
- 10 ms seek time
- 100 MB/s IO

$$|DEPT| \equiv 100 \equiv 1 \text{ page} \equiv 10 \text{ KB}$$

$$|EMP| \equiv 10 \text{ K} \equiv 100 \text{ pages} \equiv 1 \text{ MB}$$

$$|KIDS| \equiv 30 \text{ K} \equiv 300 \text{ pages} \equiv 3 \text{ MB}$$

QUERY: select \* from EMP, DEPT, KIDS  
where e.sal > 10K AND  
EMP.dno = DEPT.dno AND  
EMP.eid = KIDS.eid;



CPU cost in terms of predicate evaluations.

Note we need to know the physical op. implementation

What about I/O cost?

(11)

- Assume DEPT is outer in the NL join.

↳ 1 scan of DEPT (1 page) : 10ms seek

↳ 100 sequential scans of EMP (100 × 100 pages) [Only 10 pages fit in cache, so we need to read data from disk]

↳ 1 scan of EMP : 1 seek + read in 1MB

↓  
100 pages × 10KB/page ≈ 1MB

↳ 10ms (seek) +  $\frac{1\text{MB}}{100\text{MB/s}} = 20\text{ms}$ .

↳ In total 20ms × 100 scans = 2s.

↳ 2.01s (2s EMP + 0.01s DEPT)

- Assume DEPT is inner in the NL join

↳ read page of EMP (seek) : 10ms

↳ read DEPT into RAM : 10ms

↳ seek back to EMP : 10ms (pointer was at DEPT)

↳ scan EMP : 10ms ( $\frac{1\text{MB}}{100\text{MB/s}}$ )

↳ Total cost of 40ms

• How to choose the appropriate op. implementation. How to avoid I/O costs?

## Access Methods:

↳ Strategies to access data with minimum I/O cost.

↳ Indexes. Built to be external, i.e., to operate on disk.

↳ Different indexes have different properties (hash index not good for range queries).

↳ Indexes: Additional structure that avoids scanning all tuples. We want to use indexes when they're available.

↳ General idea of indexes:

↳ insert (key, record-id) // points from a key to a record id.

↳ lookup (key) // returns record id.

↳ lookup (low key ... high key) // return records.

↳ It'd be simple to design these in-memory, but for full generality we must support them on disk.

→ Heap Scan: we've seen this. Iterate over tuples over pages.

→ Hash File

→ B-Tree

## Hash File:

- $\text{map}(\text{Key}) \rightarrow \text{rfd}$  ; Hash table that for EHP hashes on 'name' attribute.
- $\rightarrow h(\text{name}) : [1, K]$   
 $\rightarrow h(x) = x \bmod K$
- Suppose  $K$  buckets, and one page per bucket.
  - $\rightarrow$  When inserting a tuple, we hash on the attribute to determine in which page to store it, i.e., we append the record/tuple to that page.
- When we receive a query that asks for "Tim", we hash the name and know in which page to find him.
  - $\rightarrow$  As opposed to scanning all pages.
- The key challenge is in how to select the # buckets.
  - $\rightarrow$  we could choose a high number  $\rightarrow$  this may be wasteful.
  - $\rightarrow$  if we choose too few, then buckets overflow
    - $\rightarrow$  we can create chains of pages, but then we progressively lose the index benefit

## + Extensible hashing:

- Create family  $H_k(x)$  of hash functions parametrized by 'k'.
  - $\rightarrow K=1$  ;  $h_1(x) = \{0, 1\}$   $h_k(x) = x \bmod 2^k$
  - $\rightarrow K=2$  ;  $h_2(x) = \{0, 1, 2, 3\}$
- Start with the  $h_k(x)$  hash function. When bucket overflows, redistribute that bucket into the new buckets given by the next hash function.

## B+Tree:

- A balanced tree in which internal nodes direct the search for some point/range query and the leaf nodes contain/point to the pages with the data.
- Internal nodes designed so they fit (each one) in one page. That way root of the tree fits in memory.
  - $\rightarrow$  Underlying records are sorted and those pages are linked to each other
  - $\rightarrow$  Link pointers / Doubly-linked list. Useful to answer range queries efficiently.
  - $\rightarrow$  Cost of traversing tree from root to leaf and then reaching the necessary pages.
- Why do we want the tree to be balanced?
- High fan-out (children per node) so height remains low. Why do we want low height?