

# 6.830 Lecture 11

Recap


10/15/2018

# Celebration of Knowledge

- 1.5h
- No phones, No laptops
- Bring your Student-ID
- The 5 things allowed on your desk
  - Calculator allowed
  - 4 pages (2 pages double sided) of your liking
  - Student-id
  - The exam (will have extra pages for notes)
  - Sugary stuff (i.e., food, drinks to keep you going)

# ER Model


# Clicker: ER Model



Suppose that there are 300 authors, 1000 books and 10 editors. What is the maximal number of triples that R contains?

- a) 300
- b) 3000
- c) 10,000
- d) 300,000
- e) 3,000,000

# Clicker: ER Model



Suppose that there are 300 authors, 1000 books and 10 editors. What is the maximal number of triples that R contains?

- a) 300
- b) 3000
- c) 10,000
- d) 300,000
- e) 3,000,000

**Solution:  $300 * 1000 = 300,000$**

**As every author and book unique identifies the editor**

# Type of Indexes

# Question

What indexes would you create for the following queries (assuming each query is the only query the database runs)

```
SELECT MAX(sal) FROM emp
```

B+Tree on emp.sal

```
SELECT sal FROM emp WHERE id = 1
```

Hash index on emp.id

```
SELECT name FROM emp WHERE sal > 100k
```

B+Tree on emp.sal (maybe)

```
SELECT name FROM emp WHERE sal > 100k AND dept = 2
```

B+tree on emp.sal (maybe), Hash on dept.dno (maybe)

# Clicker

What is the estimated IO cost for

```
SELECT *  
FROM T  
WHERE T.time = 109808098
```

Assuming that id is unique and is indexed using a non-clustered B-Tree with fanout B and N records

- a) 1      b) 2      c) 3      d)  $\text{LOG}_B N$       e)  $\text{LOG}_B N + 1$



# Clicker - Solution

You can argue that probably most of those answers are correct:

- a) If you assume the index is in memory and you only have a random IO for the data
- b & c) Most B-Trees in practice are only 2-3 levels deep with any reasonable fanout. Assuming that the root-note is in memory (a very reasonable assumption), it requires 1-2 random IOs to traverse the remaining memory, plus one random IO to get to the data page.
- d) This is the cost to traverse the B-Tree. However, if you assume that the root note is in memory as before, this is a good general estimate
- e) Is the general case and implicitly assume even the root node is not in memory.

Important: It depends on the assumptions and you should always state them. For Query Optimizers Constants matter a lot (O-Notation is not enough).

# Joins

# Clicker: Index Nested Loop

```
for s in S
    find matches in R
```

What is the expect IO cost assuming a non-clustered hash-index?

- a)  $|S| + \{S\}$
- b)  $|S| + \{R\}$
- c)  $|S| + \{S\} * 2$

$|S|$  = NB of pages,  $\{S\}$  = NB of records

# Clicker: Solution

```
for s in S
    find matches in R
```

What is the expected IO cost assuming a non-clustered hash-index?

a)  $|S| + \{S\}$  The most common (simplified) cost estimate

b)  $|S| + \{R\}$  Usually better if  $\{R\} > \{S\}$

c)  $|S| + \{S\} * 2$  Considers the indirection for a non-clustered index.

Question: What about  $|S| + \{R\} * 2$ . When is that the better estimate?

# Clicker Question

- Assuming disk can do 100 MB/sec I/O, and 10ms / seek
- And the following schema:

```
grades (cid int, g_sid int, grade char(2))  
students (s_int, name char(100))
```

1. Estimate time to sequentially scan grades, assuming it contains 1M records (Consider: field sizes, headers)
2. Estimate time to join these two tables, using nested loops, assuming students fits in memory but grades does not, and students contains 10K records.

# Clicker Question

- Assuming disk can do 100 MB/sec I/O, and 10ms / seek
- And the following schema:

```
grades (cid int, g_sid int, grade char(2))  
students (s_int, name char(100))
```

1. Estimate time to sequentially scan grades, assuming it contains 1M records (Consider: field sizes, headers)

(a) .23 sec    (b) 0.251 sec    (c) .21 sec    (d) 0.5 sec

Int = 8 Bytes, 1MB = 1000KB

# Seq Scan Grades

```
grades (cid int, g_sid int, grade char(2))
```

Record-Size:

8 bytes (cid) + 8 bytes (g\_sid) + 2 bytes  
(grade) + 4 bytes (header) = **22 bytes**

Data-Size: **22 MB**

Scan-Time: 22 MB / 100 MB/sec + seek-time  
= .22 sec + 0.01 sec ➔ **.23 sec**

# Clicker Question

Assuming disk can do 100 MB/sec I/O, and 10ms / seek, and the following schema:

```
grades (cid int, g_sid int, grade char(2))  
students (s_int, name char(100))
```

**Estimate time to join these two tables, using nested loops, assuming students fits in memory but grades does not, and students contains 10K records.**

(a) .241 sec   (b) 0.251 sec   (c) .211 sec   (d) 0.502 sec

Int = 8 Bytes, 1MB = 1000KB



# Grace Hash

Algorithm:

Partition:

Suppose we have  $P$  partitions, and  $H(x) \rightarrow [0 \dots P-1]$

Choose  $P = |S| / M \rightarrow P \leq \sqrt{|S|}$  //may need to leave a little slop for imperfect hashing

Allocate  $P$  1-page output buffers, and  $P$  output files for  $R$

For each  $r$  in  $R$ :

Write  $r$  into buffer  $H(r)$

If buffer full, append to file  $H(r)$

Allocate  $P$  output files for  $S$

For each  $s$  in  $S$ :

Write  $s$  into buffer  $H(s)$

if buffer full, append to file  $H(s)$

*Need one page of RAM for each of  $P$  partitions*

*Since*

*$M > \sqrt{|S|}$  and*

*$P \leq \sqrt{|S|}$ , all is well*

Join:

For  $i$  in  $[0, \dots, P-1]$

Read file  $i$  of  $R$ , build hash table

Scan file  $i$  of  $S$ , probing into hash table and outputting matches

Total I/O cost: Read  $|R|$  and  $|S|$  twice, write once  
 **$3(|R| + |S|)$  I/Os**

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
|    |    |    |
|    |    |    |

P output buffers

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

P output files

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
|    |    | 5  |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
|    | 4  | 5  |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 3  | 4  | 5  |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  |    |    |

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

| R0 | R1 | R2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  |    |    |

Need to flush R0 to F0!

| F0 | F1 | F2 |
|----|----|----|
|    |    |    |
|    |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
|    | 4  | 5  |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
| 3  |    |    |
| 6  |    |    |
|    |    |    |
|    |    |    |



# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 4  | 5  |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
| 3  |    |    |
| 6  |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 4  | 5  |
|    |    | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  |    |    |
| 6  |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 4  | 5  |
|    | 1  | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  |    |    |
| 6  |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 4  | 5  |
|    | 1  | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  |    |    |
| 6  |    |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  |    | 5  |
|    |    | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  |    |
| 6  | 1  |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 7  | 5  |
|    |    | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  |    |
| 6  | 1  |    |
|    |    |    |
|    |    |    |

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

| R0 | R1 | R2 |
|----|----|----|
| 9  | 7  | 5  |
|    |    | 14 |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  |    |
| 6  | 1  |    |
|    |    |    |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 7  |    |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
|    |    |    |
|    |    |    |



# Example

$$P = 3; H(x) = x \bmod P$$



R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

| R0 | R1 | R2 |
|----|----|----|
| 9  | 7  | 11 |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
|    |    |    |
|    |    |    |

# Example

$P = 3; H(x) = x \bmod P$



$R = 5, 4, 3, 6, 9, 14, 1, 7, 11$

$S = 2, 3, 7, 12, 9, 8, 4, 15, 6$

| R0 | R1 | R2 |
|----|----|----|
|    |    |    |
|    |    |    |

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11



S=2,3,7,12,9,8,4,15,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:

R=5,4,3,6,9,14,1,7,11


S=2,3,7,12,9,8,4,15,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files



| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:  
3,3

R=5,4,3,6,9,14,1,7,11


S=2,3,7,12,9,8,4,15,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files



| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

Matches:  
3,3

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

Scan F0 from S



# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

Load F0 from R into memory

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |



Scan F0 from S

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

7,7

4,4

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

# Example

$$P = 3; H(x) = x \bmod P$$

R=5,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

Matches:

3,3

9,9

6,6

7,7

4,4

R Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 4  | 5  |
| 6  | 1  | 14 |
| 9  | 7  | 11 |
|    |    |    |

S Files

| F0 | F1 | F2 |
|----|----|----|
| 3  | 7  | 2  |
| 12 | 4  | 8  |
| 9  |    |    |
| 15 |    |    |
| 6  |    |    |

# Sort-Merge Join

$(|R| + |S| < M \text{ -- in memory})$

Sort R, Sort S

Merge

$|R| + |S|$  i/os

# Join Processing in Database Systems with Large Main Memories

LEONARD D. SHAPIRO  
North Dakota State University

# Sort Merge Join

Equi-join of two tables S & R

$|S|$  = Pages in S;  $\{S\}$  = Tuples in S

$|S| \geq |R|$

M pages of memory;  $M > \sqrt{|S|}$

Algorithm:

- Partition S and R into memory sized sorted runs, write out to disk
- Merge all runs simultaneously

Total I/O cost: Read  $|R|$  and  $|S|$  twice, write once

**$3(|R| + |S|)$  I/Os**



# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

If each run is M pages and  $M > \sqrt{|S|}$ , then there are at most

$$R1 = \frac{|S|}{\sqrt{|S|}} = \sqrt{|S|}$$

S1 = runs of S

So if  $|R| = |S|$ , we actually need M to be  $2 \times \sqrt{|S|}$

[handwavy argument in paper for why it's only  $\sqrt{|S|}$ ]

| R |    |    |   |    |    |
|---|----|----|---|----|----|
| 1 |    |    |   |    |    |
| 3 |    |    |   |    |    |
| 4 | 14 | 11 | 7 | 12 | 15 |

OUTPUT

Need enough memory to keep 1 page of each run in memory at a time

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|-----|-----|-----|-----|-----|-----|
| 1   | 6 ← | 1   | 2 ← | 8 ← | 4 ← |
| 3 ← | 9   | 7 ← | 3   | 9   | 6   |
| 4   | 14  | 11  | 7   | 12  | 15  |

| OUTPUT |
|--------|
|        |
|        |
|        |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|-----|-----|-----|-----|-----|-----|
| 1   | 6 ← | 1   | 2   | 8 ← | 4 ← |
| 3 ← | 9   | 7 ← | 3 ← | 9   | 6   |
| 4   | 14  | 11  | 7   | 12  | 15  |

| OUTPUT |
|--------|
|        |
|        |
|        |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|-----|-----|-----|-----|-----|-----|
| 1   | 6 ← | 1   | 2   | 8 ← | 4 ← |
| 3 ← | 9   | 7 ← | 3 ← | 9   | 6   |
| 4   | 14  | 11  | 7   | 12  | 15  |

| OUTPUT |
|--------|
| (3,3)  |
|        |
|        |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|-----|-----|-----|-----|-----|-----|
| 1   | 6 ← | 1   | 2   | 8 ← | 4 ← |
| 3   | 9   | 7 ← | 3 ← | 9   | 6   |
| 4 ← | 14  | 11  | 7   | 12  | 15  |

| OUTPUT |
|--------|
| (3,3)  |
| (4,4)  |
|        |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4

R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|-----|-----|-----|-----|-----|-----|
| 1   | 6 ← | 1   | 2   | 8 ← | 4 ← |
| 3   | 9   | 7 ← | 3   | 9   | 6   |
| 4 ← | 14  | 11  | 7 ← | 12  | 15  |

| OUTPUT |
|--------|
| (3,3)  |
| (4,4)  |
|        |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|---|---|---|---|---|---|
| 1   | 6  | 1   | 2   | 8  | 4   |
| 3   | 9   | 7  | 3   | 9   | 6  |
| 4  | 14  | 11  | 7  | 12  | 15  |

| OUTPUT |
|--------|
| (3,3)  |
| (4,4)  |
| (6,6)  |
|        |

# Example

R=1,4,3,6,9,14,1,7,11

S=2,3,7,12,9,8,4,15,6

R1 = 1,3,4







R2 = 6,9,14

R3 = 1,7,11

S1 = 2,3,7

S2 = 8,9,12

S3 = 4,6,15

| R1  | R2  | R3  | S1  | S2  | S3  |
|---|---|---|---|---|---|
| 1   | 6  | 1   | 2   | 8  | 4   |
| 3   | 9   | 7  | 3   | 9   | 6  |
| 4  | 14  | 11  | 7  | 12  | 15  |

...

| OUTPUT |
|--------|
| (3,3)  |
| (4,4)  |
| (6,6)  |
| (7,7)  |

Output in  
sorted  
order!



# Summary

Notation:  $P$  partitions / passes over data; assuming hash is  $O(1)$

| Sort-Merge   | Grace Hash                                     |
|--|--|
| I/O: $3( R  +  S )$<br>CPU: $O(P \times \{S\}/P \log \{S\}/P)$ | I/O: $3( R  +  S )$<br>CPU: $O(\{R\} + \{S\})$ |

# Summary

Notation: P partitions / passes over data; assuming hash is  $O(1)$

| Sort-Merge  | Grace Hash                                      |
|---|---|
| I/O: $3 ( R  +  S )$<br>CPU: $O(P \times \{S\}/P \log \{S\}/P)$ | I/O: $3 ( R  +  S )$<br>CPU: $O(\{R\} + \{S\})$ |

Grace hash is generally a safe bet, unless memory is close to size of tables

Extra cost of sorting makes sort merge unattractive unless there is a way to access tables in sorted order (e.g., a clustered index), or a need to output data in sorted order (e.g., for a subsequent ORDER BY)

# Clicker

## Assume a Star-Schema:

```
fact(o_id, p_id, zip, discount, amount, )  
product(p_id, price, description)  
location(zip, city, s_id)
```


## And the following query:

```
SELECT *  
FROM fact f , product p, location l  
WHERE f.zip = l.zip AND f.p_id = p.p_id  
AND city=Boston
```

Fact contains 10M, Product contains 10k, Location contains 500  
1k records fit into main memory

**Would a reasonable optimizer pick the plan  
on the right?**

- a) Yes (and why)
- b) No (and why not)



# Clicker


## Assume a Star-Schema:

```
fact(o_id, p_id, zip, discount, amount, )  
product(p_id, price, description)  
location(zip, city, s_id)
```

## And the following query:

```
SELECT *  
FROM fact f , product p, location l  
WHERE f.zip = l.zip AND f.p_id = p.p_id  
AND city=Boston
```

Fact contains 10M, Product contains 10k, Location contains 500  
1k records fit into main memory



**Solution: No, you wouldn't use a grace hash join; after the city-filter it would fit into main memory – Hash Join or Nested Loop Join are better options. Furthermore, you would not use a nested loop join for p, as products don't fit into main memory**